

Sitong Fang

✉ sitongfang1@gmail.com · sitongfang.com ·  [Sitong Fang](#)

Research Interests: Trustworthy multimodal AI, world models that truly understand physics, and AI alignment.

Education

Peking University

Yuanpei College, Artificial Intelligence

B.S., 2023 – 2027 (expected)

- **Honors:** SenseTime Scholarship; Yuanpei Young Scholar; Boya Scholarship; Soong Ching Ling Future Scholarship; Beijing Natural Science Foundation *QiYan* Research Program; Freshman Scholarship (First Prize)
- **Technical Skills:** Python, PyTorch, vLLM, DeepSpeed, Ray, CUDA, C/C++, $\text{\LaTeX}$

Publications

When Slower Isn't Truer: Inverse Scaling Law of Truthfulness in Multimodal Reasoning ACL 2026

Sitong Fang, Wenjing Cao, Jiahao Li, Xuyao Wang, Chi-Min Chan, Sirui Han, Juntao Dai, Yike Guo, Yaodong Yang, Jiaming Ji

- Discovered an inverse scaling law: slower reasoning models are less truthful in multimodal settings.
- Proposed **TruthfulVQA**, the first benchmark for multimodal truthfulness evaluation with 5,000+ images, and **TruthfulJudge**, a human-in-the-loop evaluation framework.

Debate with Images: Detecting Deceptive Behaviors in Multimodal LLMs ICML 2026

Sitong Fang, Shiyi Hou, Kaile Wang, Boyuan Chen, Donghai Hong, Jiayi Zhou, Juntao Dai, Yaodong Yang, Jiaming Ji

- Introduced **MM-DeceptionBench**, the first benchmark for evaluating deceptive behaviors in multimodal LLMs.
- Proposed **Debate with Images**, a multi-agent debate framework requiring models to ground claims in visual evidence, significantly improving deception detection accuracy.

AI Deception: Risks, Dynamics, and Controls

Under Review 

Boyuan Chen*, Sitong Fang*, Jiaming Ji*, ..., Yaodong Yang†, Tiejun Huang†, Ya-Qin Zhang†, HongJiang Zhang†, Andrew Yao†

* Equal contribution † Corresponding author

- The first systematic international report on AI deception, with Turing Award laureate Andrew Yao as the corresponding author. Submitted to ACM Computing Surveys.
- Formally defined AI deception using signaling theory, analyzed the deception cycle, and proposed mitigation strategies.

Research Experience

Physis AI

Research Scientist

Feb 2026 – Present

- World model startup building world foundation models that truly understand physics. Backed by Hillhouse Ventures and Yanyuan Capital (over \$10M first round).

PKU-Alignment Group, Peking University

Research Intern · Advised by [Prof. Yaodong Yang](#)

Dec 2024 – Present

- Led the development of **TruthfulVQA**, the first benchmark for multimodal truthfulness evaluation (ACL 2026).
- Led the development of **Debate with Images** (ICML 2026), a visually grounded multi-agent debate framework for multimodal deception detection, along with **MM-DeceptionBench**, the first benchmark for multimodal deception.
- Co-first authored the **AI Deception Survey**, with Turing Award laureate Andrew Yao as corresponding author.
- Core contributor to [Align-Anything](#) (4,600+ stars on GitHub), an all-modality alignment framework.

- Core contributor to [Eval-Anything](#), a comprehensive safety evaluation framework for multimodal models.

HKGAI R&D Center / HKUST
Research Intern

Dec 2024 – Mar 2025

- Contributed to **HKGAI-V1**, the Hong Kong government’s first locally fine-tuned generative AI model based on DeepSeek, supporting Cantonese, Mandarin, and English with localized safety alignment.

Honors and Awards

2026	SenseTime Scholarship	<i>30 recipients nationwide</i>
2025	Yuanpei Young Scholar, Peking University	<i>10 annual recipients</i>
2025	Beijing Natural Science Foundation Undergraduate <i>QiYan</i> Research Program	<i>Sole recipient in cohort</i>
2025	Soong Ching Ling Future Scholarship	<i>10 annual recipients university-wide</i>
2024	Boya Scholarship, Peking University	
2024	CMSC Scholarship, Peking University	
2024	Academic Excellence Award, Peking University	
2024	Social Service Award, Peking University	
2023	Freshman Scholarship (First Prize), Peking University	
2023	Ranked 1st in Fujian Province, National College Entrance Exam (Science)	